DS311 - R Lab Assignment

Ana Paula Felix de Queiroz

2023-04-14

R Assignment 1

- In this assignment, we are going to apply some of the build in data set in R for descriptive statistics analysis.
- To earn full grade in this assignment, students need to complete the coding tasks for each question to get the result.
- After finished all the questions, knit the document into HTML format for submission.

Question 1

Using the **mtcars** data set in R, please answer the following questions.

```
# Loading the data
data(mtcars)

# Head of the data set
head(mtcars)
```

```
##                   mpg cyl disp  hp drat    wt  qsec vs am gear carb
## Mazda RX4         21.0   6  160 110 3.90 2.620 16.46  0  1    4    4
## Mazda RX4 Wag     21.0   6  160 110 3.90 2.875 17.02  0  1    4    4
## Datsun 710        22.8   4  108  93 3.85 2.320 18.61  1  1    4    1
## Hornet 4 Drive    21.4   6  258 110 3.08 3.215 19.44  1  0    3    1
## Hornet Sportabout 18.7   8  360 175 3.15 3.440 17.02  0  0    3    2
## Valiant           18.1   6  225 105 2.76 3.460 20.22  1  0    3    1
```

a. Report the number of variables and observations in the data set.

```
# Enter your code here!
dim(mtcars)
```

```
## [1] 32 11
```

```
# Answer:
print("There are total of 11 variables and 32 observations in this data set.")
```

```
## [1] "There are total of 11 variables and 32 observations in this data set."
```

b. Print the summary statistics of the data set and report how many discrete and continuous variables are in the data set.

```
# Enter your code here!

summary(mtcars)
```

```
##       mpg             cyl             disp             hp
##  Min.   :10.40   Min.   :4.000   Min.   : 71.1   Min.   : 52.0
##  1st Qu.:15.43   1st Qu.:4.000   1st Qu.:120.8   1st Qu.: 96.5
##  Median :19.20   Median :6.000   Median :196.3   Median :123.0
##  Mean   :20.09   Mean   :6.188   Mean   :230.7   Mean   :146.7
##  3rd Qu.:22.80   3rd Qu.:8.000   3rd Qu.:326.0   3rd Qu.:180.0
##  Max.   :33.90   Max.   :8.000   Max.   :472.0   Max.   :335.0
##       drat             wt             qsec             vs
##  Min.   :2.760   Min.   :1.513   Min.   :14.50   Min.   :0.0000
##  1st Qu.:3.080   1st Qu.:2.581   1st Qu.:16.89   1st Qu.:0.0000
##  Median :3.695   Median :3.325   Median :17.71   Median :0.0000
##  Mean   :3.597   Mean   :3.217   Mean   :17.85   Mean   :0.4375
##  3rd Qu.:3.920   3rd Qu.:3.610   3rd Qu.:18.90   3rd Qu.:1.0000
##  Max.   :4.930   Max.   :5.424   Max.   :22.90   Max.   :1.0000
##        am              gear            carb
##  Min.   :0.0000   Min.   :3.000   Min.   :1.000
##  1st Qu.:0.0000   1st Qu.:3.000   1st Qu.:2.000
##  Median :0.0000   Median :4.000   Median :2.000
##  Mean   :0.4062   Mean   :3.688   Mean   :2.812
##  3rd Qu.:1.0000   3rd Qu.:4.000   3rd Qu.:4.000
##  Max.   :1.0000   Max.   :5.000   Max.   :8.000
```

```
continuous <- sum(sapply(mtcars, is.double))
discrete <- ncol(mtcars) - continuous
continuous
```

```
## [1] 11
```

```
discrete
```

```
## [1] 0
```

```
# Answer:
print("There are 0 discrete variables and 11 continuous variables in this data set.")
```

```
## [1] "There are 0 discrete variables and 11 continuous variables in this data set."
```

c. Calculate the mean, variance, and standard deviation for the variable **mpg** and assign them into variable names m, v, and s. Report the results in the print statement.

```
# Enter your code here!
attach(mtcars)
m=mean(mpg)
v=sd(mpg)^2
s = sd(mpg)

print(paste("The average of Mile Per Gallon from this data set is ",m  , " with variance ",v  , " and s
```

```
## [1] "The average of Mile Per Gallon from this data set is  20.090625  with variance  36.3241028225806
```

d. Create two tables to summarize 1) average mpg for each cylinder class and 2) the standard deviation of mpg for each gear class.

```
# Enter your code here!

library(dplyr)
```

```
##
## Attaching package: 'dplyr'

## The following objects are masked from 'package:stats':
##
##     filter, lag

## The following objects are masked from 'package:base':
##
##     intersect, setdiff, setequal, union
```

```
# Create a table of average mpg for each cylinder class
mpg_by_cyl <- mtcars %>%
  group_by(cyl) %>%
  summarise(avg_mpg = mean(mpg))

mpg_by_cyl
```

```
## # A tibble: 3 x 2
##     cyl avg_mpg
##   <dbl>   <dbl>
## 1     4    26.7
## 2     6    19.7
## 3     8    15.1
```

```
# Create a table of standard deviation of mpg for each gear class
mpg_by_gear <- mtcars %>%
  group_by(gear) %>%
  summarise(sd_mpg = sd(mpg))

# Print the table
print(mpg_by_gear)
```

```
## # A tibble: 3 x 2
##    gear sd_mpg
##   <dbl>  <dbl>
## 1     3   3.37
## 2     4   5.28
## 3     5   6.66
```

e. Create a crosstab that shows the number of observations belong to each cylinder and gear class combinations. The table should show how many observations given the car has 4 cylinders with 3 gears, 4 cylinders with 4 gears, etc. Report which combination is recorded in this data set and how many observations for this type of car.

```r
# Enter your code here!
crosstab = table(mtcars$cyl,mtcars$gear,
                 dnn = c("Cyl","gears"))
crosstab
```

```
##    gears
## Cyl  3  4  5
##   4  1  8  2
##   6  2  4  1
##   8 12  0  2
```

```r
print("The most common car type in this data set is car with 8 cylinders and 3 gears. There are total o
```

```
## [1] "The most common car type in this data set is car with 8 cylinders and 3 gears. There are total
```

Question 2

Use different visualization tools to summarize the data sets in this question.

a. Using the **PlantGrowth** data set, visualize and compare the weight of the plant in the three separated group. Give labels to the title, x-axis, and y-axis on the graph. Write a paragraph to summarize your findings.

```r
# Load the data set
data("PlantGrowth")

# Head of the data set
head(PlantGrowth)
```

```
##   weight group
## 1   4.17  ctrl
## 2   5.58  ctrl
## 3   5.18  ctrl
## 4   6.11  ctrl
## 5   4.50  ctrl
## 6   4.61  ctrl
```

```r
# Enter your code here!

library(ggplot2)
```

```
##
## Attaching package: 'ggplot2'
```

```
## The following object is masked from 'mtcars':
##
##     mpg
```

```
ggplot(PlantGrowth, aes(x=group, y=weight, fill=group)) +
  geom_boxplot() + labs(title = "Weight of Plants by Group", x = "Group", y = "Weight")
```

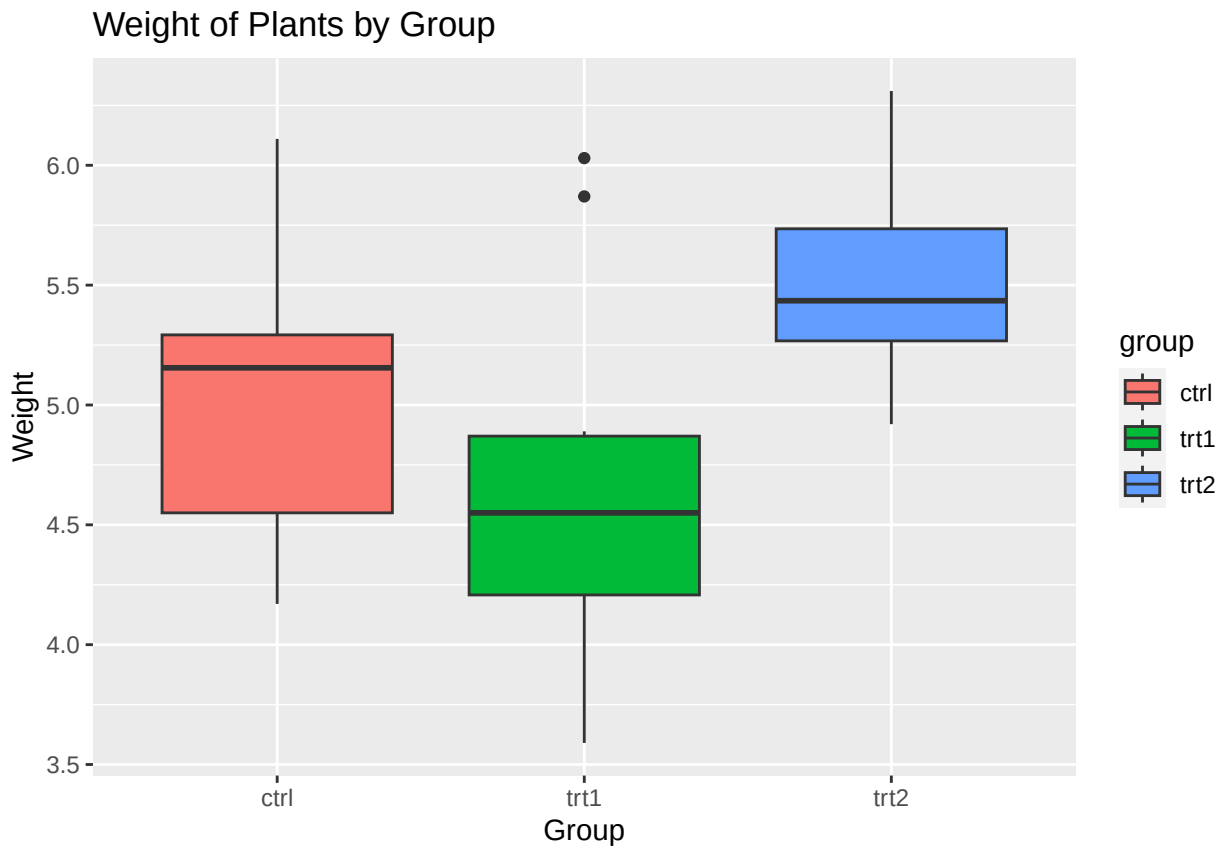

Result:

=> Report a paragraph to summarize your findings from the plot!

The box plot shows that weights of plants for group TRT2 is higher than the other two with a median at 5.4. While the group Trt1 has the lower median around 4.6. Group CTRL seems to have the largest spread of the 50% of the data. The distribution of group ctrl seem to be skewed to the write while group trt2 seem to be skewed to the left.We are also able to identify that trt1 has outliers.

b. Using the **mtcars** data set, plot the histogram for the column **mpg** with 10 breaks. Give labels to the title, x-axis, and y-axis on the graph. Report the most observed mpg class from the data set.

```
hist(mtcars$mpg, breaks = 10,
     main = "MPG distribution",
     xlab = "MPG",
     ylab = "Frequency")
```

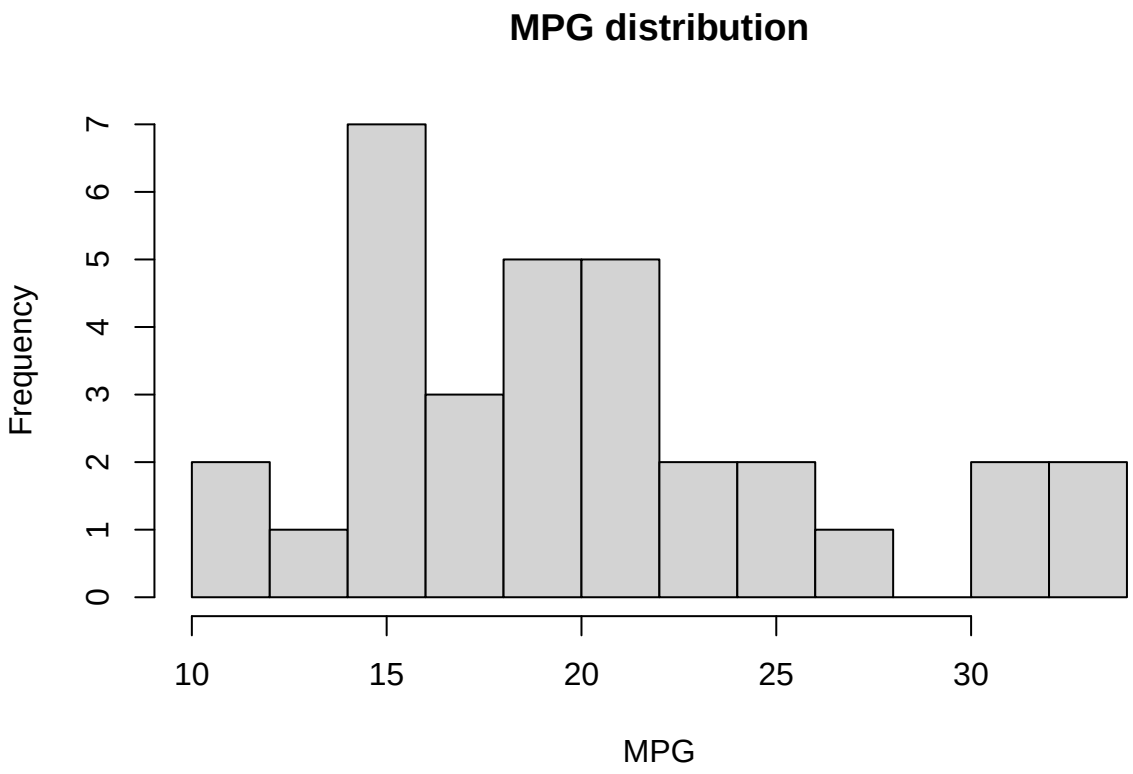


```
print("Most of the cars in this data set are in the class of 10 mile per gallon.")
```

```
## [1] "Most of the cars in this data set are in the class of 10 mile per gallon."
```

c. Using the **USArrests** data set, create a pairs plot to display the correlations between the variables in the data set. Plot the scatter plot with **Murder** and **Assault**. Give labels to the title, x-axis, and y-axis on the graph. Write a paragraph to summarize your results from both plots.

```
# Load the data set
data("USArrests")

# Head of the data set
head(USArrests)
```

```
##            Murder Assault UrbanPop Rape
## Alabama      13.2     236       58 21.2
## Alaska       10.0     263       48 44.5
## Arizona       8.1     294       80 31.0
## Arkansas      8.8     190       50 19.5
## California    9.0     276       91 40.6
## Colorado      7.9     204       78 38.7
```

```r
# Enter your code here!


pairs(USArrests, main = "Pairs Plot")
```

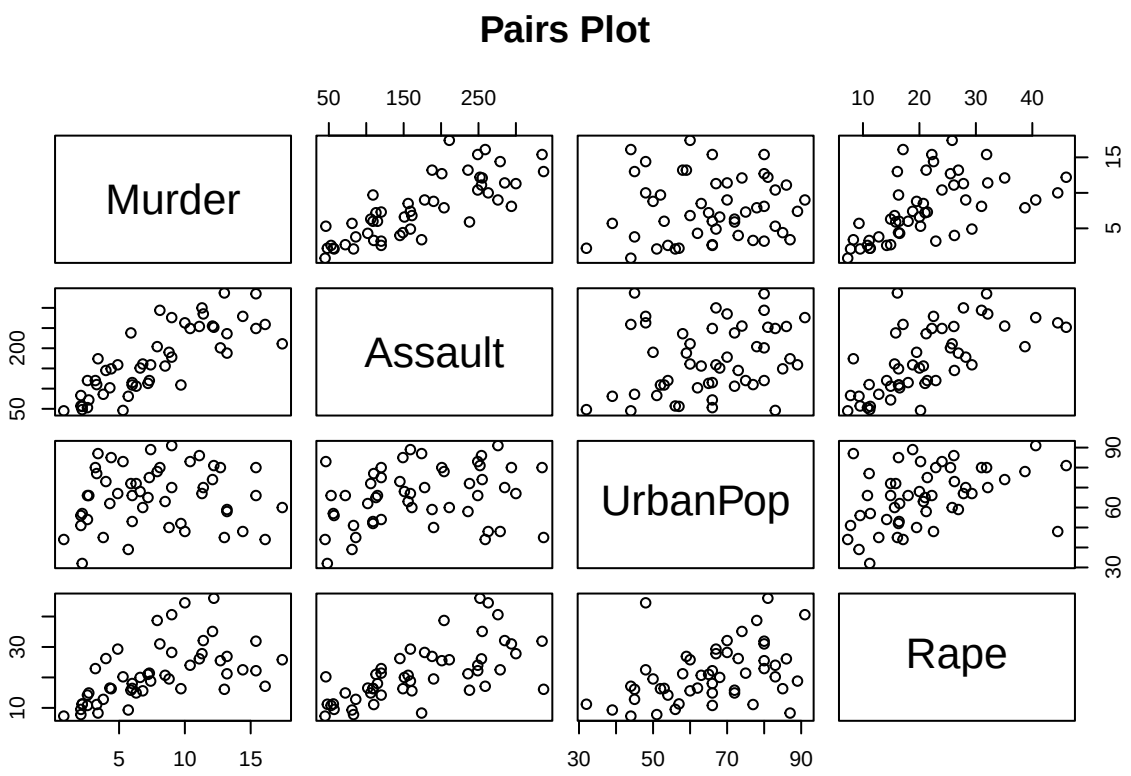


```r
plot(USArrests$Murder, USArrests$Assault,
     main = "Scatter Plot of Murder and Assault",
     xlab = "Murder", ylab = "Assault")
```

Result:

=> Report a paragraph to summarize your findings from the plot! It seems like murder and Assault have the highest positive correlation. We could say that there is a moderate positive correlation between Rape and murder as well. It seems like Urban Population is the variable with the lowest correlation with Assault and murder.

Question 3

Download the housing data set from www.jaredlander.com and find out what explains the housing prices in New York City.

Note: Check your working directory to make sure that you can download the data into the data folder.

```
## Warning in dir.create("data"): 'data' already exists
```

a. Create your own descriptive statistics and aggregation tables to summarize the data set and find any meaningful results between different variables in the data set.

```
# Head of the cleaned data set
head(housingData)
```

```
##   Neighborhood Market.Value.per.SqFt      Boro Year.Built
```

```
## 1    FINANCIAL               200.00 Manhattan       1920
## 2    FINANCIAL               242.76 Manhattan       1985
## 4    FINANCIAL               271.23 Manhattan       1930
## 5      TRIBECA               247.48 Manhattan       1985
## 6      TRIBECA               191.37 Manhattan       1986
## 7      TRIBECA               211.53 Manhattan       1985
```

```
unique(housingData$Boro)
```

```
## [1] "Manhattan"     "Brooklyn"      "Queens"        "Bronx"
## [5] "Staten Island"
```

```
summary(housingData)
```

```
##  Neighborhood       Market.Value.per.SqFt     Boro             Year.Built
##  Length:2530        Min.   : 10.66        Length:2530        Min.   :1825
##  Class :character   1st Qu.: 75.10        Class :character   1st Qu.:1926
##  Mode  :character   Median :114.89        Mode  :character   Median :1986
##                     Mean   :133.17                           Mean   :1967
##                     3rd Qu.:189.91                           3rd Qu.:2005
##                     Max.   :399.38                           Max.   :2010
```

```
# Enter your code here!
#table discribing mean and median grouped by value by var Boro
library(dplyr)
housingData %>%
  group_by(Boro) %>%
  summarize(mean_value = mean(Market.Value.per.SqFt),
            median_value = median(Market.Value.per.SqFt))
```

```
## # A tibble: 5 x 3
##   Boro          mean_value median_value
##   <chr>              <dbl>        <dbl>
## 1 Bronx               47.9         47.4
## 2 Brooklyn            80.1         81.6
## 3 Manhattan          181.         184.
## 4 Queens              77.4         66.9
## 5 Staten Island       41.3         41.0
```

```
housingData %>%
  group_by(Neighborhood) %>%
  summarize(mean_value = mean(Market.Value.per.SqFt),
            median_value = median(Market.Value.per.SqFt)) %>%
  arrange(desc(mean_value))
```

```
## # A tibble: 148 x 3
##    Neighborhood            mean_value median_value
##    <chr>                        <dbl>        <dbl>
##  1 MIDTOWN CBD                   234.         227.
##  2 FLATIRON                      223.         230.
##  3 MIDTOWN WEST                  222.         223.
```

```
##  4 UPPER EAST SIDE (59-79)       217.         218.
##  5 CHELSEA                       216.         214.
##  6 MIDTOWN EAST                  211.         220.
##  7 EAST VILLAGE                  207.         200.
##  8 MURRAY HILL                   206.         209.
##  9 UPPER EAST SIDE (79-96)       202.         210.
## 10 GREENWICH VILLAGE-WEST        202.         214.
## # ... with 138 more rows
```

```
summary(housingData$Market.Value.per.SqFt)
```

```
##    Min. 1st Qu.  Median    Mean 3rd Qu.    Max.
##   10.66   75.10  114.89  133.17  189.91  399.38
```

b. Create multiple plots to demonstrates the correlations between different variables. Remember to label all axes and give title to each graph.

```
# Enter your code here!
library(ggplot2)
ggplot(housingData, aes(x = Boro, y = Market.Value.per.SqFt)) +
  geom_point() +
  labs(title = "Market Value by Boro",
       x = "Boro",
       y = "Market Value per SqFt")
```

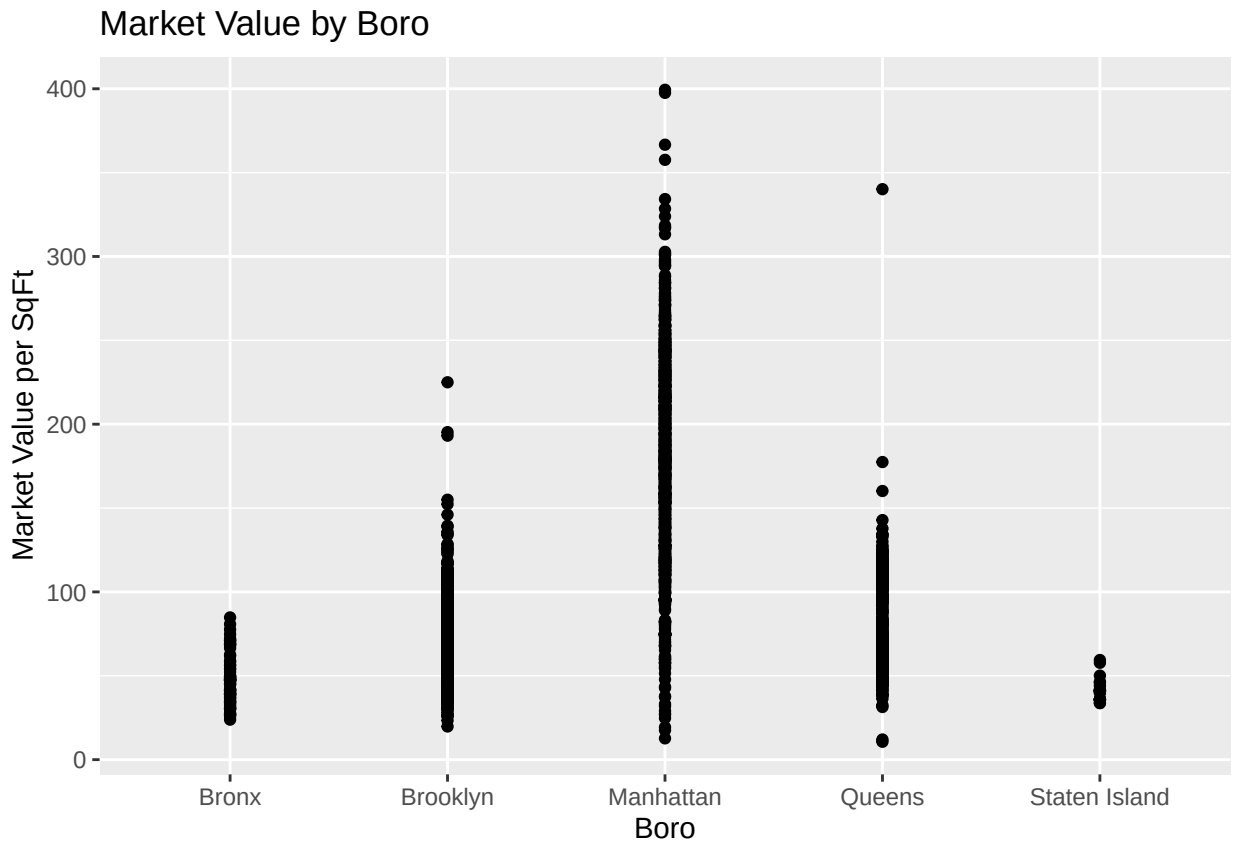

```r
top10_neighborhoods <- housingData %>%
  group_by(Neighborhood) %>%
  summarize(mean_value = mean(Market.Value.per.SqFt)) %>%
  arrange(desc(mean_value)) %>%
  top_n(5, mean_value)


ggplot(top10_neighborhoods, aes(x = reorder(Neighborhood, mean_value), y = mean_value)) +
  geom_bar(stat = "identity", fill = "steelblue") +
  labs(title = "Top 5 Neighborhoods by Market Value per SqFt",
       x = "Neighborhood", y = "Mean Market Value per SqFt")
```

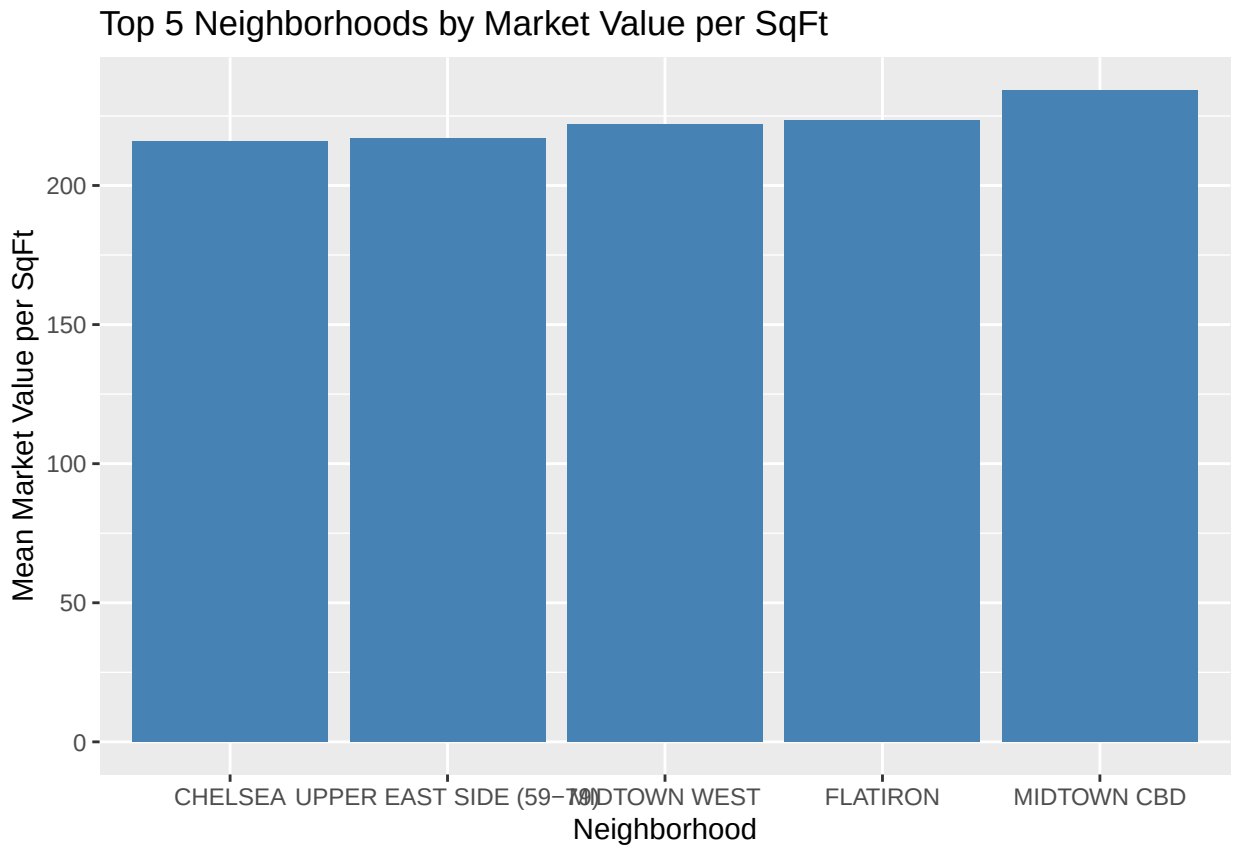


```r
ggplot(housingData, aes(x = Year.Built, y = Boro)) +
  geom_boxplot() +
  labs(title = "Distribution of Boro by Year Built",
       x = "Year Built",
       y = "Boro")
```

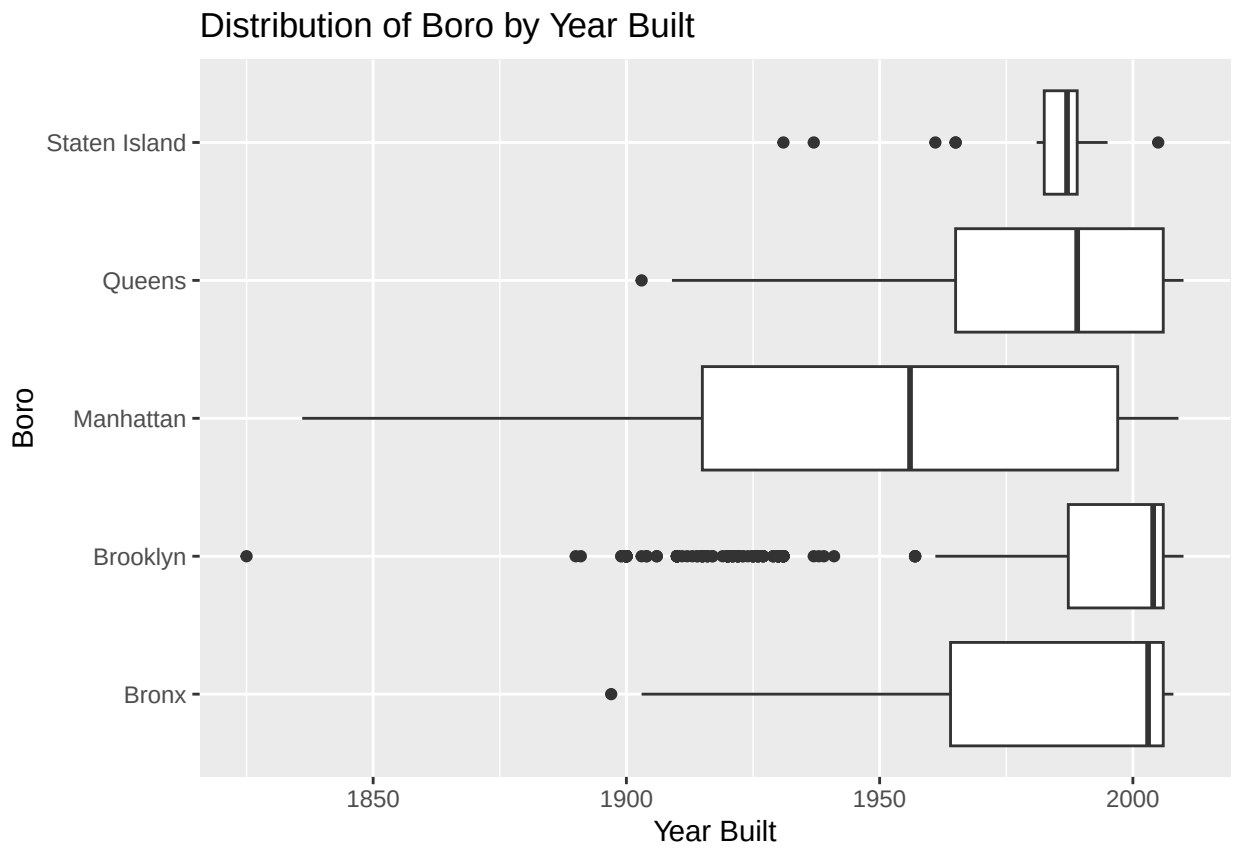

c. Write a summary about your findings from this exercise.

If we look at the variables Boro and market value we can see that Mahattan has the highest market value per square ft, queens and brooklyn seem to have a similar Market price exept for outliers. its followed by Bronx and then Staten Island.

Among all the Neighboors Midtown, Flatiron, Midtown West , uper East Side and Chelsea hank on the top highest mean for Market value per square ft

We see that broklyn has one observation as the oldest building, and manhatan has the largest distribution of buildings age with the median around 1960.